

# Latent-space Physics: Towards Learning the Temporal Evolution of Fluid Flow

Steffen Wiewel, Moritz Becher, Nils Thuerey  
Technical University of Munich

# Motivation

- The field of computational methods has been highly successful at developing powerful **numerical algorithms** to predict how the natural phenomena under consideration will behave.
- Take a different view:
  - instead of relying on analytic expressions
  - use a deep learning approach to infer physical functions based on data
  - focus on temporal evolution of complex functions in the context of fluid flows

Challenge: high dimensionality of Eulerian space-time data sets

# Contributions

- a first LSTM architecture to predict temporal evolutions of dense, 3D functions of physics system in latent spaces
- a neural-network(an encoder and decoder architecture) based simulation algorithm with significant practical speed-ups. (a strong compression)
  - more than two orders of magnitudes faster than a traditional pressure solver
- a detailed evaluation of training modalities
- Give NNs predictive capabilities for complex inverse problems

# Background: flow physics

- *Navier-Stokes* (NS) model

$\mathbf{u}$ : flow velocity

$$\partial \mathbf{u} / \partial t + \mathbf{u} \cdot \nabla \mathbf{u} = -1/\rho \nabla p + \nu \nabla^2 \mathbf{u} + \mathbf{g}, \quad \nabla \cdot \mathbf{u} = 0, \quad (1)$$

$p$ : pressure

$\rho$ ,  $\nu$ ,  $\mathbf{g}$ : density, kinematic viscosity and external forces

# Method

- **Goal:** predict future states of a physical function  $\mathbf{x}$  ( $\mathbb{R}^3 \times \mathbb{R}^+ \rightarrow \mathbb{R}^d$ ):
  - $d = 1$  for scalar functions such as pressure
  - $d = 3$  for vectorial functions such velocity fields
  - Can represent a spatio-temporal pressure function, or a velocity field
- **Key:** reducing the dimensionality of the problem using *convolutional neural networks* (CNNs) with respect to both **time and space**
- **Process:**
  - Map the original 3D problem into a smaller spatial *latent space*; Learn the inverse mapping at the same time
  - Maps a collection of reduced representations into an encoding of the temporal evolution
  - Reduced temporal state is used to output a sequence of spatial latent space representations
  - Decode latent space representations to yield the full spatial data set

# Method (cont'd)

- Goal: Given a functional representation  $f_t$ , a current state and  $n$  previous states, predict future states from  $\mathbf{x}(t+h)$  to  $\mathbf{x}(t+oh)$ :  $f_t(\mathbf{x}(t-nh), \dots, \mathbf{x}(t-h), \mathbf{x}(t)) \approx [\mathbf{x}(t+h), \dots, \mathbf{x}(t+oh)]$ . (2)

- Employ  $f_d$  (CNN / decoder) and  $f_e$  (CNN / encoder) to compute low dimensional encodings  
Rephrase Eq. (2):

$$\begin{aligned}\tilde{f}_t(f_e(\mathbf{x}(t-nh)), \dots, f_e(\mathbf{x}(t))) &\approx [\mathbf{c}^{t+h}, \dots, \mathbf{c}^{t+oh}] \\ f_d([\mathbf{c}^{t+h}, \dots, \mathbf{c}^{t+oh}]) &\approx [\mathbf{x}(t+h), \dots, \mathbf{x}(t+oh)] \quad (3) \\ f_d(\tilde{f}_t(f_e(\mathbf{x}(t-nh)), \dots, f_e(\mathbf{x}(t)))) &\approx f_t(\mathbf{x}(t-nh), \dots, \mathbf{x}(t))\end{aligned}$$

- $\mathbf{c}$ : latent space,  $\mathbf{c} \in \mathbb{R}^{m_s}$
- $\tilde{f}_t$ : the prediction network, also an encoder-decoder structure
  - turns the temporal stack of encoded data points  $f_e(\mathbf{x})$  into a reduced representation  $\mathbf{d}$ , refer to as *temporal context*.

# Architecture (spatial network)

- a fully convolutional autoencoder

$$\min_{\theta_d, \theta_e} |f_d(f_e(\mathbf{x}(t))) - \mathbf{x}(t)|_2,$$

$\theta_d, \theta_e$ : parameters of the decoder and encoder

- Note that: paths from  $f_{e_k}$  over  $c_k$  to  $f_{e_5}$  are only active in pretraining stage k. After pretraining only the path  $f_{d_k}$  over  $c_5$  to  $f_{d_5}$  remains active.

Goal: Reducing Dimensionality

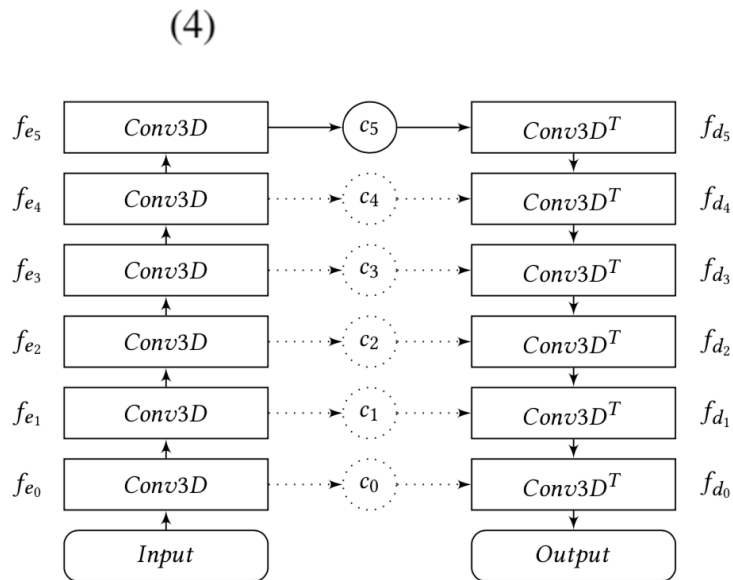


Fig 1: Overview of the autoencoder network

# Architecture (temporal network)

Goal: Prediction of Future States

- LSTM layers for predicting the evolution over time.
  - Minimize the mean absolute error between the o predicted and ground truth states:

$$\min_{\theta_t} \left| \tilde{f}_t(\mathbf{c}^{t-nh}, \dots, \mathbf{c}^{t-h}, \mathbf{c}^t) - [\mathbf{c}^{t+h}, \dots, \mathbf{c}^{t+oh}] \right|_1. \quad (6)$$

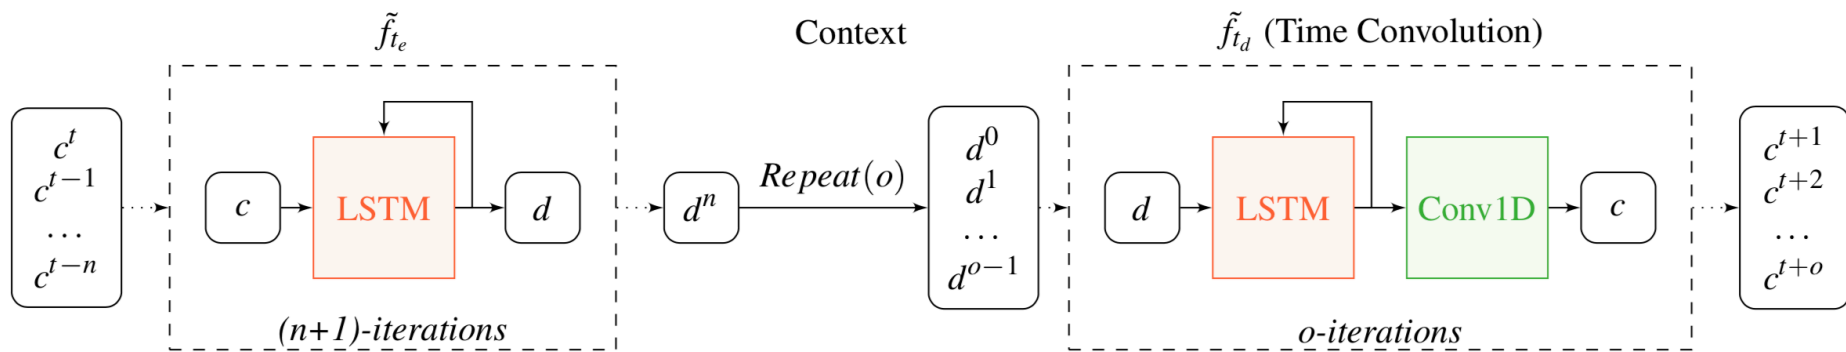
- a recurrent encoder module (LSTM)
    - transform  $n + 1$  latent space points into a temporal context  $\mathbf{d}$
  - a recurrent decoder module with similar structure (layers of LSTM units)
    - takes a context  $\mathbf{d}$  as input, outputs a series of future, spatial latent space representations
- ✓ During encoding, only keep the very last context, pass it to the decoder part.
- ✓ During decoding, context is repeated o times for the decoder LSTM to infer the desired future states.



# Architecture (temporal network)

Goal: Prediction of Future States

- To prevent overfitting: a hybrid structure of LSTM units and convolutions



dashed boxes indicate an  
iterative evaluation

Navier-Stokes (NS) model:

$$\partial \mathbf{u} / \partial t + \mathbf{u} \cdot \nabla \mathbf{u} = -1/\rho \nabla p + \nu \nabla^2 \mathbf{u} + \mathbf{g}, \quad \nabla \cdot \mathbf{u} = 0,$$

# Datasets: Fluid Flow Data

- Rely on a NS solver with operator splitting to calculate the training data at discrete points in space and time
  - Computing motion of the fluid with the help of transport, i.e. *advection* steps, for the velocity  $\mathbf{u}$ , evaluating external forces  $\mathbf{g}$ , then computing the harmonic pressure function  $p$ .
  - A visible, passive quantity such as smoke density  $\rho$ , or a level-set representation  $\phi$  for free surface flows is advected in parallel to the velocity itself.
  - Consider a *split* pressure. a hydrostatic pressure  $p_s$  for cell at height  $z$
  - a hydrostatic pressure  $p_s$  for cell at height  $z$ :
$$p_s(z) = p(z_0) + \frac{1}{A} \int_{z_0}^z \iint_A \mathbf{g} \rho(h) \, dx dy \, dh$$
with  $z_0$ ,  $p_0$ ,  $A$  denote surface height, surface pressure, and cell area.
  - Density and gravity can be treated as constant in this setting, simplifies to:  $p_s = \rho \mathbf{g}(z - z_0)$ 
    - Can be evaluated efficiently
    - Drawback: only valid for fluids in hydrostatic equilibrium, typically cannot be used for dynamic simulations in 3D


# Datasets: Fluid Flow Data (cont'd)

- Given a data-driven method to predict pressure fields, we can incorporate the hydrostatic pressure into a 3D liquid simulation:
  - decomposing the regular pressure field  $p_t$  into:  $p_t = p_s + p_d$
  - $p_s$ : hydrostatic component
  - $p_d$ : dynamic component
  - our autoencoder separately receives and encodes the two fields  $p_s$  and  $p_d$ .
- With split pressure, the autoencoder could potentially put more emphasis on the small-scale fluctuations  $p_d$  from the hydrostatic pressure gradient.


# Datasets:

	liquid64	liquid128	smoke128
Scenes $n_s$	4000	800	800
Time steps $n_t$	100	100	100
Size	419.43GB	671.09GB	671.09GB
Size, encoded	1.64GB	2.62GB	2.62GB

- Three 3D data sets:
- use randomized simulation setups, target scenes with high complexity:
  - strong visible splashes and vortices,
  - large CFL (Courant- Friedrichs-Lewy) numbers (how fast information travels from cell to cell in a complex simulation domain)
- For each data sets:
  1. generate  $n_s$  scenes of different initial conditions
  2. discard the first  $n_w$  time steps (typically contain small, regular, and less representative dynamics)
  3. store a fixed number of  $n_t$  time steps: a final size of  $n_s n_t$  spatial data sets.
  4. Normalize each dataset to  $[-1,1]$ .



**Figure 17:** *Examples of initial scene states in the liquid64 data set.*



**Figure 18:** *Examples of initial scene states in the liquid128 data set. The more complex initial shapes are visible in several of these configurations.*

	liquid64	liquid128	smoke128
Scenes $n_s$	4000	800	800
Time steps $n_t$	100	100	100
Size	419.43GB	671.09GB	671.09GB
Size, encoded	1.64GB	2.62GB	2.62GB



**Figure 19:** *Examples states of the smoke128 training data set at  $t = 65$ .*

# Interval Prediction

- Use network prediction for a time interval of  $i_p$  time steps
- Then perform a single full simulation step without any network calculations
  - numerical time integration + network prediction
  - Our prediction intervals  $i_p$  on the order of 4 to 14 steps
  - $i_p = 0$ : the LSTM prediction is not used at all
  - $i_p = \infty$ : the fully recurrent LSTM

# Evaluation and Training

- PSNR (peak signal-to-noise ratio) as a baseline metric
- a surface-based Hausdorff distance
- surface error: 
$$e_h = \max(1/|S_p| \sum_{\mathbf{p}_1 \in S_p} \phi_r(\mathbf{p}_1), 1/|S_r| \sum_{\mathbf{p}_2 \in S_r} \phi_p(\mathbf{p}_2))/\Delta x.$$

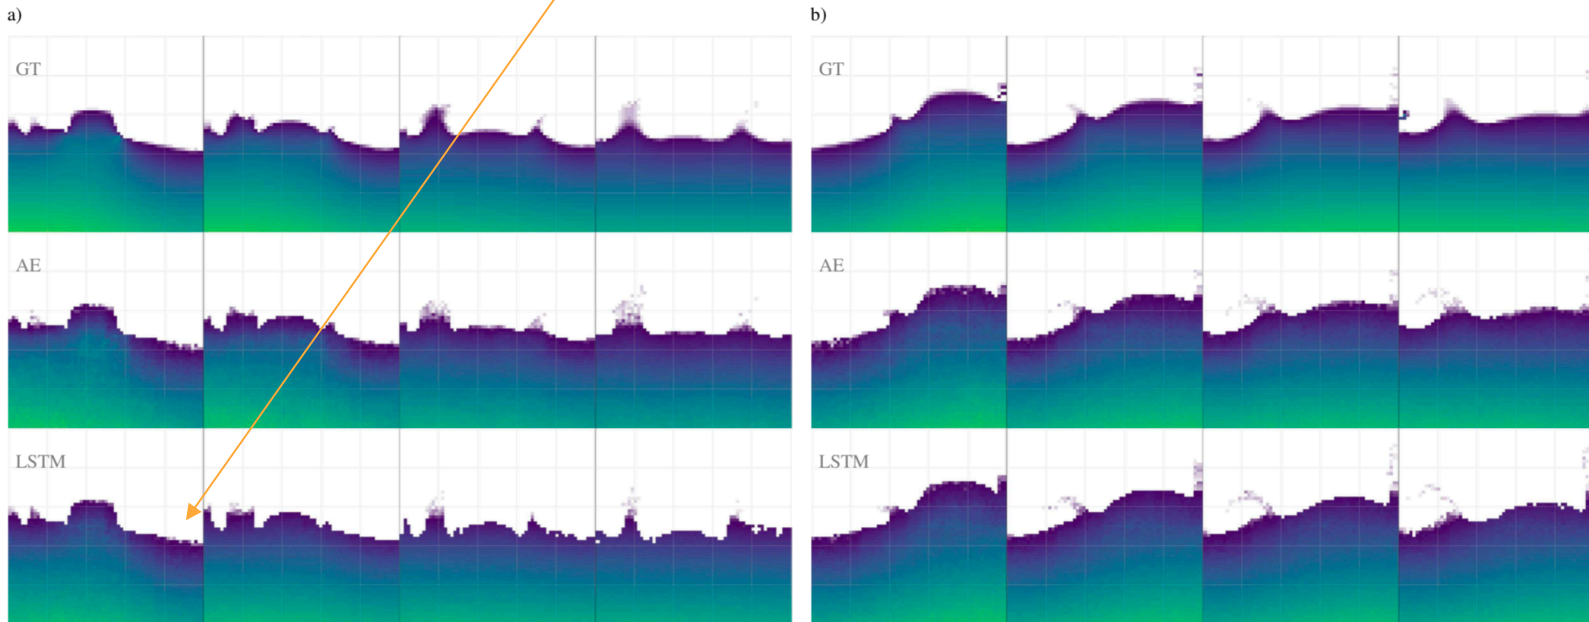
where  $\phi_r, \phi_p$  are two signed distance functions

- Spatial Encoding
  - the autoencoder network in conjunction with a numerical time integration scheme
  - encoder:  $\mathbf{c} = f_e(\mathbf{x})$ ; decoder:  $\mathbf{x}' = f_d(\mathbf{c})$
- Temporal Prediction
  - a quantity  $\mathbf{x}'$  is inferred based on a series of previous latent space points

the network successfully learned an abstraction of the temporal evolution of the flow.

# Results


- Compare ground truth, autoencoder baseline and outputs of our prediction network:
  - The temporal predictions closely match the autoencoder baseline.
  - Network can reproduce complex behavior of the underlying simulations.





# Results

- Liquid surfaces predicted by different models for 40 steps with  $i_p = \infty$ .
  - The velocity version (green) leads to large errors in surface position, all three pressure versions closely capture the large-scale motions.
  - On smaller scales, both split pressure variants ( $p_s$  and  $p_d$ ) introduce artifacts



# Results

- large speed-ups and robust simulations for a significant variety of fluid scenes

Interval $i_p$	Solve	Mean surf. dist	Speedup
Reference	2.629s	0.0	1.0
4	0.600s	0.0187	4.4
9	0.335s	0.0300	7.8
14	0.244s	0.0365	10.1
$\infty$	0.047s	0.0479	55.9

	Enc/Dec	Prediction	Speedup
Core exec. time	4.1ms + 3.3ms	9.5ms	155.6

Performance of ten liquid128 scenes, averaged over 150 simulation steps each. The mean surface distance is a measure of deviation from the reference per solve.

	Solve	Speedup
Reference	169ms	1.0


	Enc/Dec	Prediction	
Core exec., $o = 1$	3.8ms + 3.2ms	9.6ms	10.2
Core exec., $o = 3$	3.9ms + 3 * 3.3ms	12.5ms	19.3

Performance of ten liquid64 scenes, averaged over 150 simulation steps each.

# Results


- How well the model can predict future states based on a single set of inputs. For multiple output steps ( $o > 1$ ), our model predicts several latent space points from a single time context  $\mathbf{d}$ . Accuracy for 1, 3 and 5 steps of output:

- Accuracy barely degrades when multiple steps are predicted at once;
- But more efficient.  $o = 3$  prediction only requires 30% more time to evaluate, despite generating three times as many predictions



# Results

- A trained model for the *smoke128* data set.



- Despite the significantly different physics, our approach successfully predicts the evolution and motion of the vortex structures
  - However, we noticed a tendency to underestimate pressure values, and to reduce small-scale motions.

# Conclusions

- Limitations
  - LSTM strongly relies on AE (which primarily encodes large scale dynamics) while small scale dynamics are integrated by the alignment of free surface boundary conditions
  - Current simple AE can introduce a certain amount of noise; can be alleviated by different network architectures
- Conclusions
  - Neural network architectures can predict the temporal evolution of dense physical functions
  - learned latent spaces with LSTM-CNN hybrids are suitable for this task
  - Significant increases in simulation performance
    - more than 150x faster than a regular pressure solve
    - represents an important first step towards deep-learning powered simulation algorithms