

# Volumetric Isosurface Rendering with Deep Learning-Based Super-Resolution

Sebastian Weiss , Mengyu Chu, Nils Thuerey and Rüdiger Westermann  
Technical University of Munich

# Motivation

- isosurface volume ray-casting
  - A efficient rendering method to render isosurface given a viewpoint
  - Workload of ray-casting linear in the number of pixels
    - i.e. Decrease the number of pixels -> decrease workload
    - Visual quality not satisfactory → **super resolution by deep learning**.
- Consider global illumination -> increase visual quality
  - E.g. AO (ambient occlusion)
  - **Problem:** calculate AO is time-consuming
  - AO can be generated from other stored information (eg. depth and normal maps)

# Ambient occlusion

- Ambient light
  - In phong model, a fixed lighting **constant** to simulate the scattering of light
  - In reality, it should not be constant
  - Eg. darkening creases, holes and surfaces that are close to each other
    - occluded by surrounding geometry -> light rays have less places to escape -> darker



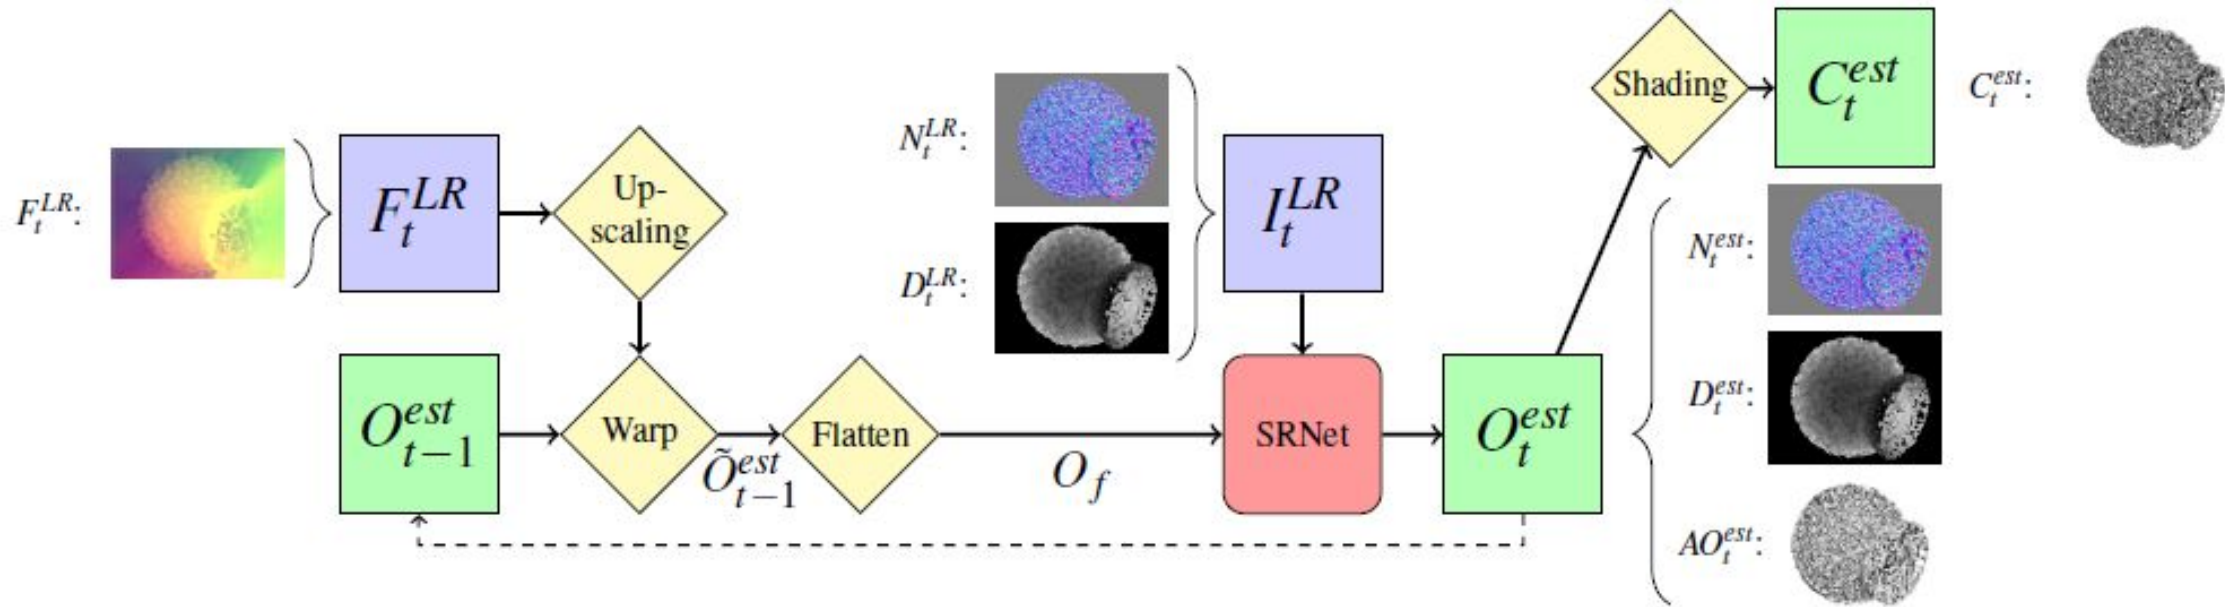
- AO: expensive
  - Reason: take surrounding geometry into account
- Paper: approximate by low-resolution stored information

# Contribution

- 1. learning the upscaling of a low resolution sampling of an isosurface -> a higher resolution
  - By reconstruction of spatial detail and shading
  - Why?
    - More information; robust to lighting
- 2. infer AO in the high resolution w/o AO in the low resolution inputs
  - No need to simulate AO
- 3. add a motion loss -> maintain frame-to-frame coherence
- 4. perform a quality evaluation

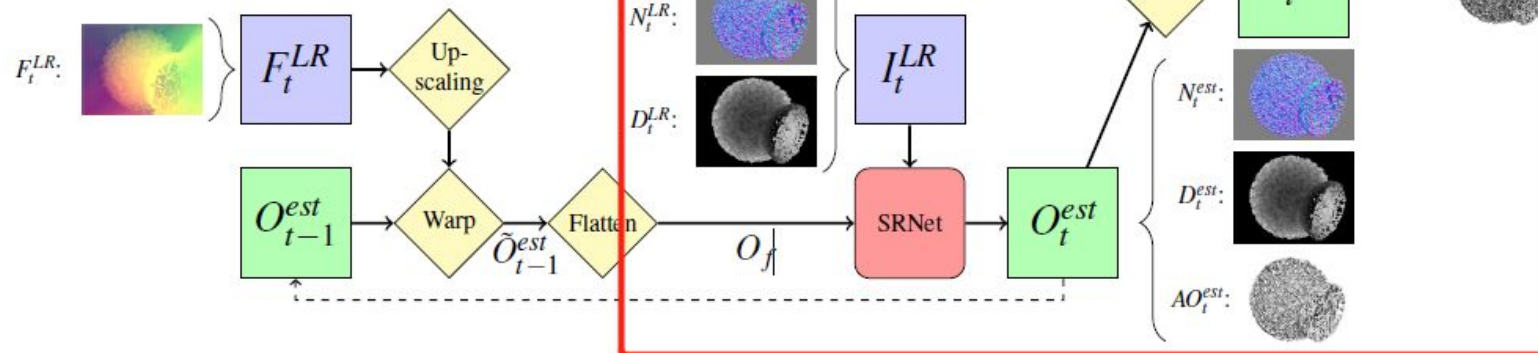
# isosurface learning

Blue: low resolution inputs green: high resolution outputs  
Yellow: fixed preprocessing steps Red: trained network



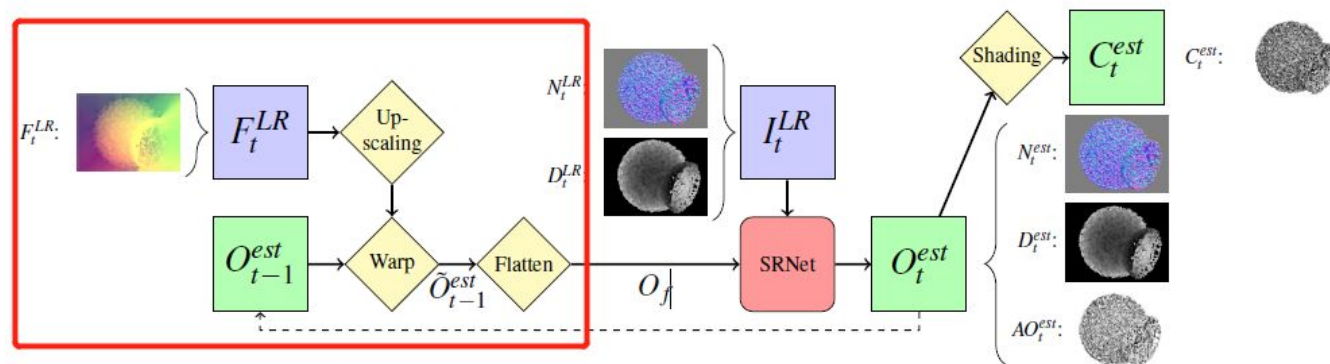
- Goal: perform 4 upscaling
  - i.e. from input images of size  $H \times W$  to output images of size  $4H \times 4W$
  - Input: LR depth and normal maps
  - Output: HR depth and normal maps, additional AO map
  - screen-space shading used for generating final color

# Input data



- Input maps generated via volumetric ray-casting
- low resolution input maps (size  $H \times W$ )
  - $M_t^{LR} \in [-1, +1]^{H \times W}$ : binary input mask
    - Whether the isosurface hit or not
    - network learn continuous value, use these to smoothly blend the final color
  - $N_t^{LR} \in [-1, +1]^{3 \times H \times W}$ : normal maps
  - $D_t^{LR} \in [0, 1]^{H \times W}$ : depth maps
  - LR input to the network  $I_t^{LR} := \{M_t^{LR}, N_t^{LR}, D_t^{LR}\} \in R^{5 \times H \times W}$

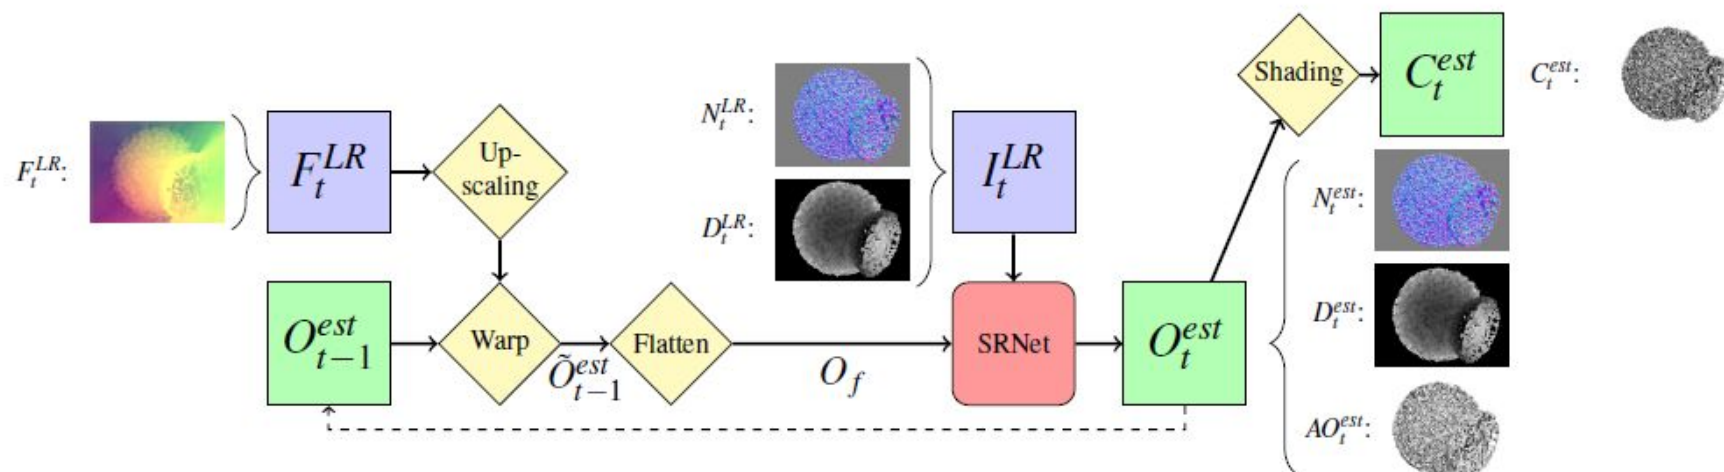
# Input data



- $F_t^{LR} \in [-1, +1]^{2 \times H \times W}$ : A map of displacement vectors
  - The screen space-flow from the previous view to current view
  - Goal: minimize the deviation of current inferred HR map from the extrapolated previous one
  - Isosurface from different viewpoints: the same coordinate at **world space**, different at **screen space**, denote as  $x'_t$  and  $x'_{t-1}$
  - Flow computed by  $f_t := x'_t - x'_{t-1}$
  - Current flow field, up-scaling via bi-linear interpolation
  - Then use inverse flow vector to determine the target location



# Input data



## • High resolution input data

- Used as GT in training process
- The same as low resolution input + AO map  $AO_t^{LR} \in [0, 1]^{4H \times 4W}$ 
  - in AO map, 0 or 1 -> no or full occlusion
- GT image
  - $O_t^{GT} := \{M_t^{GT}, N_t^{GT}, D_t^{GT}, AO_t^{GT}\} \in R^{6 \times 4H \times 4W}$

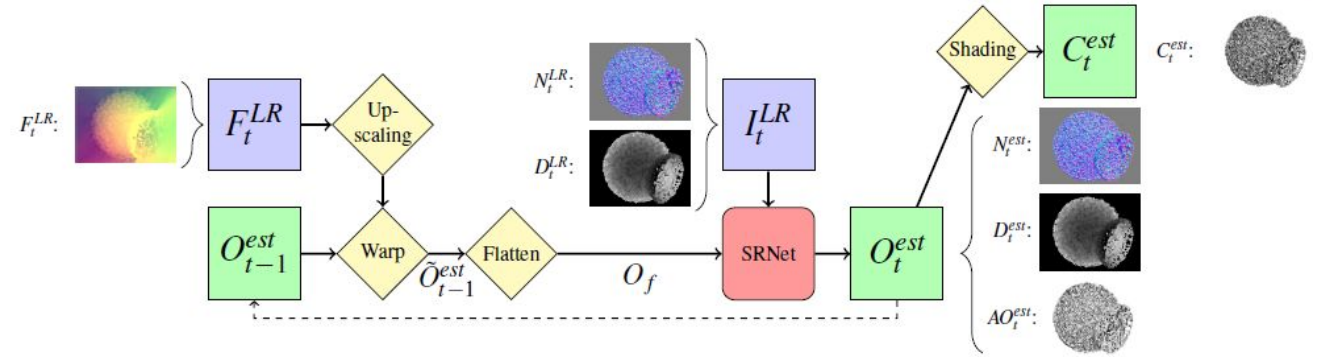


# Processing stage

- 1. upscaling & warping
- 2. flatten previous map to low resolution by apply space-to-depth transformation

$$S_s : \mathbb{R}^{6 \times 4H \times 4W} \rightarrow \mathbb{R}^{4^2 \times 6 \times H \times W}.$$

- 3. Super-Resolution:
  - Network receive  $I_t^{LR}$  (5 channels) and the flattened, warped prediction from the previous frame  $O_f$  (16 \* 6 channels), estimate the six-channel output  $O_t^{est}$ , (HR mask, normal, depth and AO maps)



- 4. Shading: generate a color image
    - screen-space Phong shading with AO
- $$C_{rgb} = \text{Phong}(c_a, c_d, c_s, c_m, N_t^{est}) * AO_t^{est},$$
- Output mask  $M_t^{est}$ 
    - Clamp first to  $[-1, +1]$ , rescale to  $[0, 1]$ , shows smooth fall-off values across edges

$$C_t^{est} = \text{lerp}(c_{bg}, C_{rgb}, M_t^{est}),$$

# Loss functions

- 1. Spatial Loss

$$\mathcal{L}_{X,L_1} = \|X_t^{est} - X_t^{GT}\|_1, \quad \mathcal{L}_{X,L_2} = \|X_t^{est} - X_t^{GT}\|_2^2.$$

- 2. Perceptual loss

- detailed outputs instead of smoothed mean values
- two images are similar <- have similar activations in the latent space of a pre-trained network

$$\mathcal{L}_{X,P} = \|\phi(X_t^{est}) - \phi(X_t^{GT})\|_2^2.$$

- VGG-19 can be used on shaded images

- Differentiable Phong shading

- Gradient can flow from loss function to update weight of neural networks

- VGG on other entries, input transformed first

- 3. temporal loss

- a temporal L2 loss
- Penalize differences between current HR and previous warped one

$$\mathcal{L}_{X,temp} = \|X_t^{est} - \tilde{X}_{t-1}^{est}\|_2^2,$$

# Loss function (cont'd)

## • 4. Loss masking

- Pixels where the mask is -1 set to the background color
- Area are masked do not contribute to the final result
- crucial step simplifies the network's task: network do not need to match masked areas

## • 5. Adversarial Training

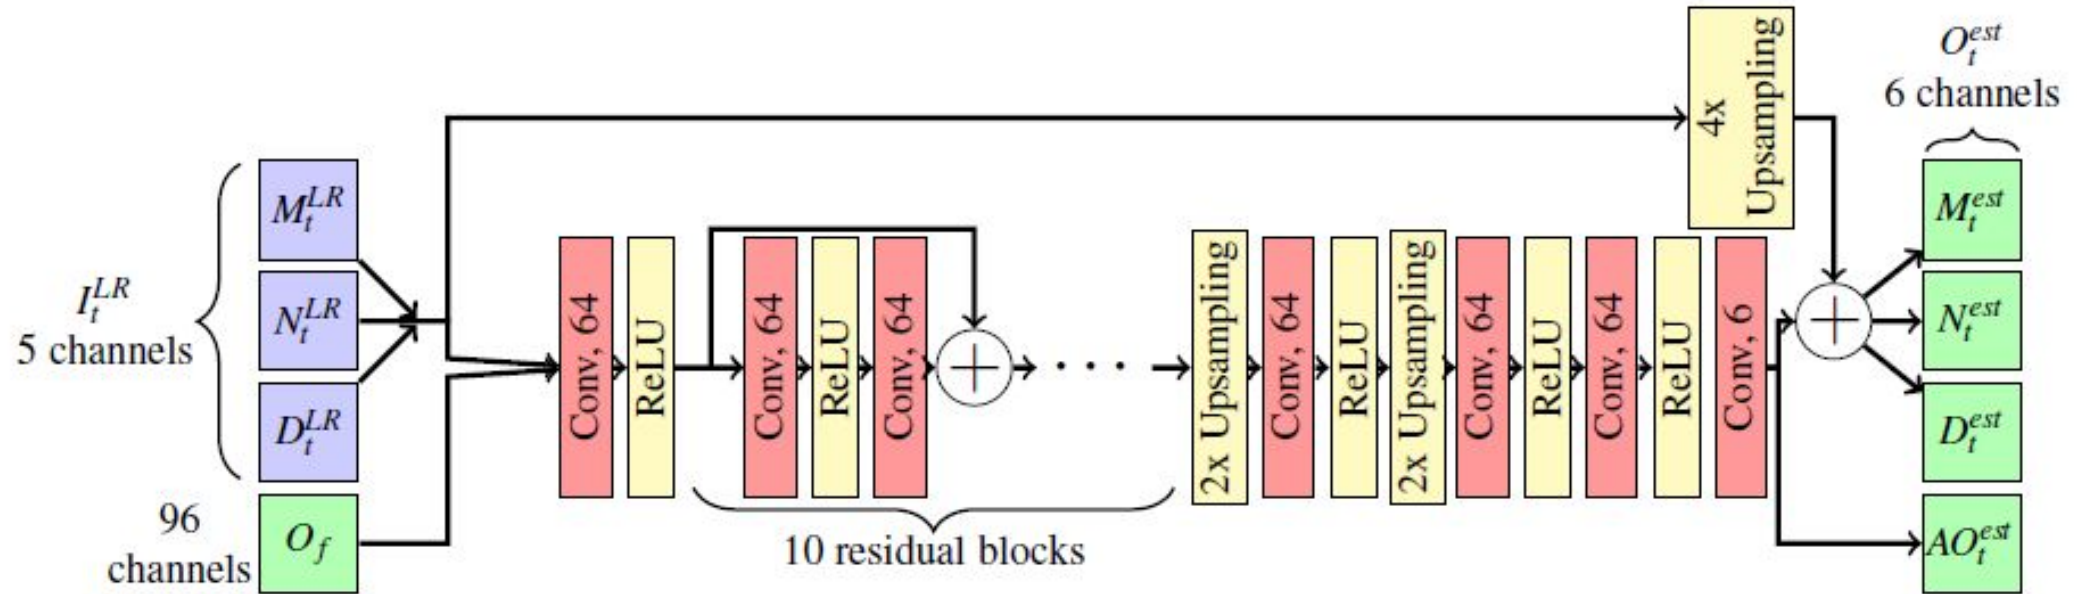
- the discriminator is provided with:
  - the high resolution output  $O_t^{est}$
  - The input  $I_t^{LR}$  as a conditional input
  - the previous frame  $I_{t-1}^{LR}$ ,  $O_{t-1}^{est}$ , to learn to penalize for temporal coherence

$$\mathcal{L}_{GAN,D} = -\log(D(x)) - \log(1 - D(G(z))).$$

$$\mathcal{L}_{GAN,G} = -\log(D(G(z)))$$

# Network architecture

- A fully convolutional frame-recurrent neural network (FRVSR-Net) consisting of a series of residual blocks



# Training data

- images of isosurfaces from different timesteps and multi-resolution versions of the Ejecta dataset.
- render 500 sequences, each consisting of 10 frames
  - Each sequence, random select start and end viewpoints
  - Construct a smooth view path
  - Get resolution 128x128
- Crop sub-regions of 32x32
  - R: benefit batch processing





Fig. 5: Example images that are used to train our networks.

- Both LR&HR image directed generated by raycaster
- Initialize previous frame of 1<sup>st</sup> frame
  - Zeroing all entries



# Loss Function Characteristics

- geometric surface properties > color images



(a)(c) super-resolution on depth and normal maps with screen-shading

(b)(d) super-resolution on color images

# Loss Function Characteristics

Network	Losses
Shaded	$\mathcal{L}_{GAN,G} + 0.5\mathcal{L}_{C,P} + 50\mathcal{L}_{C,temp}$ network acts on shaded colors
$L_1$ -color	$\mathcal{L}_{M,L_1} + \mathcal{L}_{AO,L_1} + 10\mathcal{L}_{C,L_1} + 0.1\mathcal{L}_{C,temp}$
$L_1$ -geometry	$\mathcal{L}_{M,L_1} + \mathcal{L}_{AO,L_1} + 10\mathcal{L}_{N,L_1} + 100\mathcal{L}_{D,L_1} + 0.1\mathcal{L}_{C,temp}$
Perceptual	$\mathcal{L}_{M,L_1} + \mathcal{L}_{AO,L_1} + \mathcal{L}_{N,L_1} + \mathcal{L}_{D,L_1} + 0.1\mathcal{L}_{C,temp} + 5\mathcal{L}_{N,P} + \mathcal{L}_{AO,P}$
GAN	$\mathcal{L}_{M,L_1} + \mathcal{L}_{AO,L_1} + \mathcal{L}_{N,L_1} + \mathcal{L}_{D,L_1} + 0.1\mathcal{L}_{C,temp} + \mathcal{L}_{GAN,G}$

TABLE 1: Networks and their specific loss function configurations.




Fig. 6: Visual comparison of networks with different loss function configurations: (a) Shaded, (b)  $L_1$ -color, (c)  $L_1$ -geometry (our final model), (d) Perceptual, (e) GAN, (f) ground truth



# Qualitative Evaluation

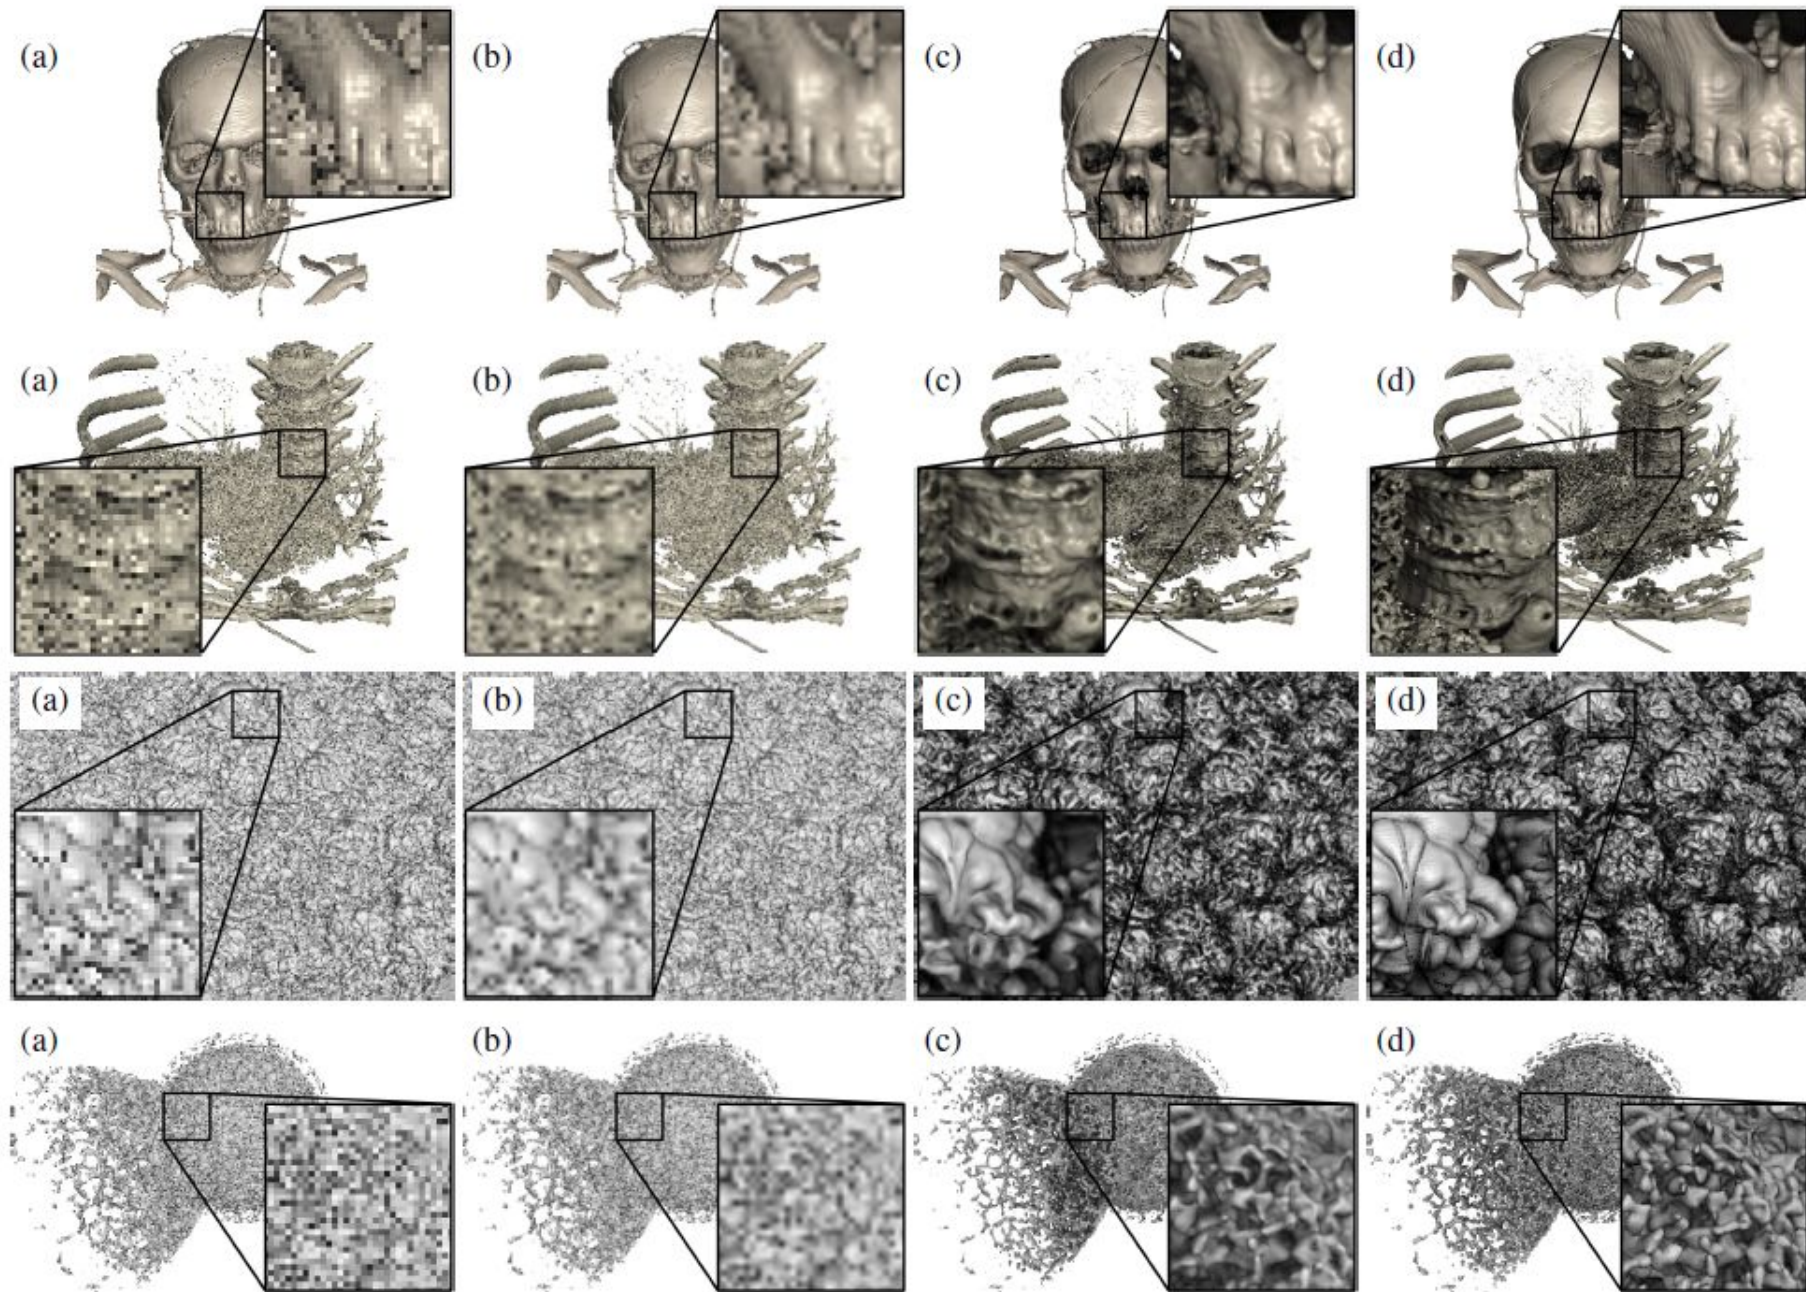


Fig. 7: Comparison of upscaling quality: (a) input, (b) bi-linear, (c) our network, (d) ground truth on the Skull, Thorax, Richtmyer-Meshkov and Ejecta dataset (top to bottom).


# Quantitative Evaluation

- peak signal-to-noise ratio (PSNR)

$$\text{PSNR}(O_t^{est}, O_t^{GT}) = -10 \log_{10}(\|O_t^{est} - O_t^{GT}\|_2^2),$$


- Structural similarity index (SSIM)

$$\text{SSIM}(O_t^{est}, O_t^{GT}) = \frac{(2\mu_{est}\mu_{GT} + c_1)(2\sigma_{est,GT} + c_2)}{(\mu_{est}^2 + \mu_{GT}^2 + c_1)(\sigma_{est}^2 + \sigma_{GT}^2 + c_2)},$$



Different loss functions


# Quantitative Evaluation (PSNR)




(a)




# Quantitative Evaluation (SSIM)




# Evaluation



(a) Network trained on w/o Clouds.



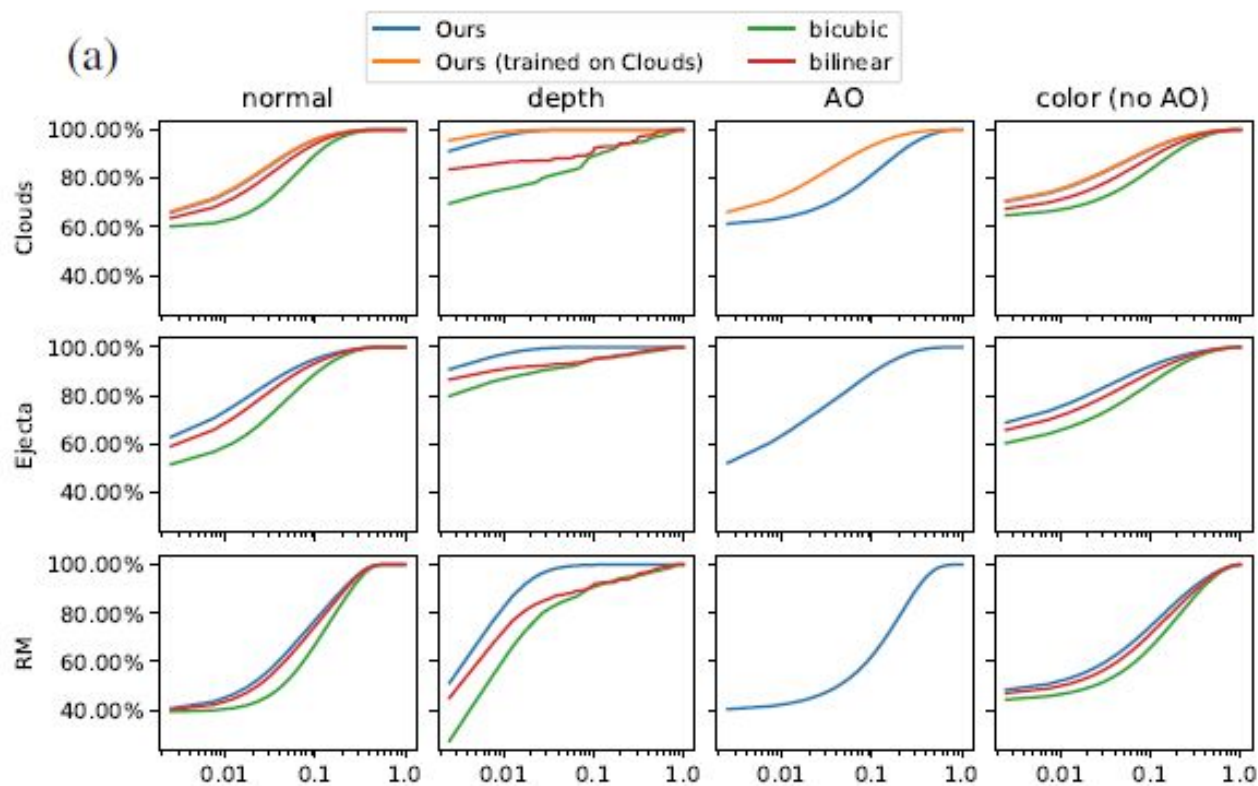
(b) Ground truth rendering of Cloud.



(c) Network re-trained only on Clouds.

# Evaluation

- Regression Error Characteristic (REC) curves
  - $REC(\tau) = P(|x^{est} - x^{GT}| \leq \tau)$




# Times for rendering

Dataset	High-res (no AO)	High-res (with AO)	Low-res	Super-res
Skull $256^3$	0.057	4.2	0.0077	0.071
Thorax $256^3$	0.069	9.1	0.010	0.071
R.-M. $1024^3$	0.088	14.5	0.014	0.072
Ejecta $1024^3$	0.163	18.6	0.031	0.072




# Application (interactive visualization)



Left: full image resolution Middle: 1/4 resolution then upscale Right: half resolution

# Application (Foveated rendering)



# Conclusion

- deep learning technique for isosurface super-resolution with AO
- For me, learning with depth/normal maps is interesting, can be applied to other generative problems too